

3D memory profile breakdown vs problem size

Best CUDA-aware FFTM p2p-waitany strategy at each N; MiB values are peak max per GPU/process.

