

4D FFTM: relative runtime by redistribution mode

Runtime is normalized by the best successful mode in each row; lower is better.

	alltoallv	alltoallw	p2p waitall	p2p waitany
slab-slab CA	2.10	1.00	1.00	1.00
slab-slab NCA	1.40	1.40	1.00	1.00
pencil-pencil CA	2.10	1.00	1.00	1.00
pencil-pencil NCA	2.23	2.23	1.00	1.00

CA = CUDA-aware MPI, NCA = non-CUDA-aware path.